

# Recitation 1

## Statistical Learning Theory

### Intro to Gradient Descent

Haoxu Huang

NYU CDS

Jan. 22, 2026

# Statistical Learning Theory

## Introduction

Section Lead TAs for this course:

- Ansh Sharma, Haoxu Huang, Patrick Shen, Yihui Hong, Karine

Office Hours:

- Ansh: Mon 11:30am - 12:30pm ET
- Haoxu: Tues 11:00am - 12:00pm ET
- Patrick: Wed 4:00pm - 5:00pm ET
- Yihui: Thurs 3:00pm - 4:00pm ET
- Karine: Fri 1:00pm - 2:00pm ET

Location:

- Room 242 in CDS (60 5th Ave, New York, NY 10011)

# Statistical Learning Theory

## Motivation

In data science, we generally need to **Make a Decision** on a problem.

To do this, we need to understand

- The setup of the problem
- The possible actions
- The effect of actions
- The evaluation of the results

How do we translate the problem into the language of DS/modeling?

# Statistical Learning Theory

## Formalization

### The Spaces

$\mathcal{X}$ : input space     $\mathcal{Y}$ : outcome space     $\mathcal{A}$ : action (decision) space

### Prediction Function

A **prediction function**  $f$  gets an input  $x \in \mathcal{X}$  and produces an action  $a \in \mathcal{A}$ :

$$f : \mathcal{X} \mapsto \mathcal{A}$$

### Loss Function

A **loss function**  $\ell(a, y)$  evaluates an action  $a \in \mathcal{A}$  in the context of an outcome  $y \in \mathcal{Y}$ :

$$\ell : \mathcal{A} \times \mathcal{Y} \mapsto \mathbb{R}$$

# Statistical Learning Theory

## Risk Function

Given a loss function  $\ell$ , how can we evaluate the “average performance” of a prediction function  $f$ ?

- To do so, we need to first assume that there is a **data generating distribution**  $P_{X,Y}$ .
- Then the expected loss of  $f$  on  $P_{X,Y}$  will reflect the notion of “average performance”.

### Definition

The **risk** of a prediction function  $f : \mathcal{X} \mapsto \mathcal{A}$  is

$$R(f) = \mathbb{E}\ell(f(x), y)$$

It is the expected loss of  $f$  on a new sample  $(x, y)$  drawn from  $P_{X,Y}$ .

# Statistical Learning Theory

Finding ‘best’ function

## Definitions

### Hypothesis Class

$\mathcal{F}$  is the family of functions we restrict our model to be. Example:  
Linear, quadratic, decision tree, two layer neural-net...

### Bayes optimal predictor within $\mathcal{F}$

$f_{\mathcal{F}}^*$  is ‘best’ function one can obtain within  $\mathcal{F}$ .

### Sample-optimal predictor

$\hat{f}_n$  is the ‘best’ function one can obtain using the data given.

### Learned predictor

$\tilde{f}_n$  is the function actually obtained using the data given.

# Statistical Learning Theory

## The Bayes Prediction Function

### Definition

A **Bayes prediction function**  $f^* : \mathcal{X} \mapsto \mathcal{Y}$  is a function that achieves the minimal risk among all possible functions:

$$f^* \in \arg \min_f R(f),$$

where the minimum is taken from all functions that maps from  $\mathcal{X}$  to  $\mathcal{A}$ .

The risk of a Bayes function is called **Bayes risk**.

# Statistical Learning Theory

## Finding the Bayes Optimal Classifier

### Goal:

In a multi-class classification problem with class  $Y \in \{0, 1, \dots, n\}$  and features  $X$ , we want to find a classifier  $h(X) = \hat{y}$  that minimizes the expected loss (risk).

$$\begin{aligned}\mathbb{E}_Y(\mathbb{I}_{\{y \neq \hat{y}\}}) &= \mathbb{E}_X \mathbb{E}_{Y|X}(\mathbb{I}_{\{y \neq \hat{y}\}} | X = x) \quad (\text{Law of total expectation}) \\ &= \mathbb{E}[\Pr(Y = 1 | X = x) \mathbb{I}_{\{\hat{y}=2,3,\dots,n\}}] \\ &\quad + \mathbb{E}[\Pr(Y = 2 | X = x) \mathbb{I}_{\{\hat{y}=1,3,\dots,n\}}] + \dots \\ &\quad + \mathbb{E}[\Pr(Y = n | X = x) \mathbb{I}_{\{\hat{y}=1,2,3,\dots,n-1\}}]\end{aligned}$$

### Note:

This is minimized by simultaneously minimizing all terms of expectation using the classifier  $h(x) = k$  with  $\operatorname{argmax}_k \Pr(Y = k | X = x)$  for each observation  $x$ .


# Statistical Learning Theory

## Finding the Bayes Optimal Classifier

- $P(Y = 1|x) + P(Y = 2|x) + \dots + P(Y = n|x) = 1$  by complement rule of probability.
- The probability of error is  $P(\text{Error}|x) = 1 - P(\text{Correct}|x)$
- To minimize the error, we need to maximize the probability of being correct.
- Consequently, for any given observation  $x$ , the strategy that minimizes risk is simply: Pick the class  $k$  that has the highest posterior probability (i.e.  $\operatorname{argmax}_k \Pr(Y = k|X = x)$ )

# Statistical Learning Theory

## Error Decomposition



# Statistical Learning Theory

## Error Decomposition

### Approximation Error

- Caused by the choice of family of functions or capacity of the model. ( $\epsilon_{approx} = R(f_{\mathcal{F}}) - R(f^*)$ )
- Expand the capacity of the model.

### Estimation Error

- Caused by finite number of data. ( $\epsilon_{est} = (R(\hat{f}_n) - R(f_{\mathcal{F}}))$ )
- Obtain more data/add regularizer

### Optimization Error

- Caused by not able to find the best parameters.  
 $(\epsilon_{opt} = R(\tilde{f}) - R(\hat{f}_n))$
- Try different optimization algorithms, learning rates, etc.

# Statistical Learning Theory

## Error Decomposition

Decomposition of “Excess Risk” (how much worse our final model is compared to the perfect Bayes optimal)

$$\begin{aligned} R(\tilde{f}_n) - R(f^*) &= \epsilon_{\text{approx}} + \epsilon_{\text{est}} + \epsilon_{\text{opt}} \\ &= (R(f_{\mathcal{F}}) - R(f^*)) \\ &\quad + (R(\hat{f}_n) - R(f_{\mathcal{F}})) \\ &\quad + (R(\tilde{f}_n) - R(\hat{f}_n)) \end{aligned}$$

# Statistical Learning Theory

## Gradient Descent

### Motivation:

Our goal is to find  $\hat{f}_n$ , the best possible model from given data

**Naive approach:** Take gradient of loss function, solve for parameters that gives you 0.

- Computationally intractable
- Impossible to compute due to complex function structure

The optimal parameters for LR:  $\hat{\beta} = (X^\top X)^{-1} X^\top Y$

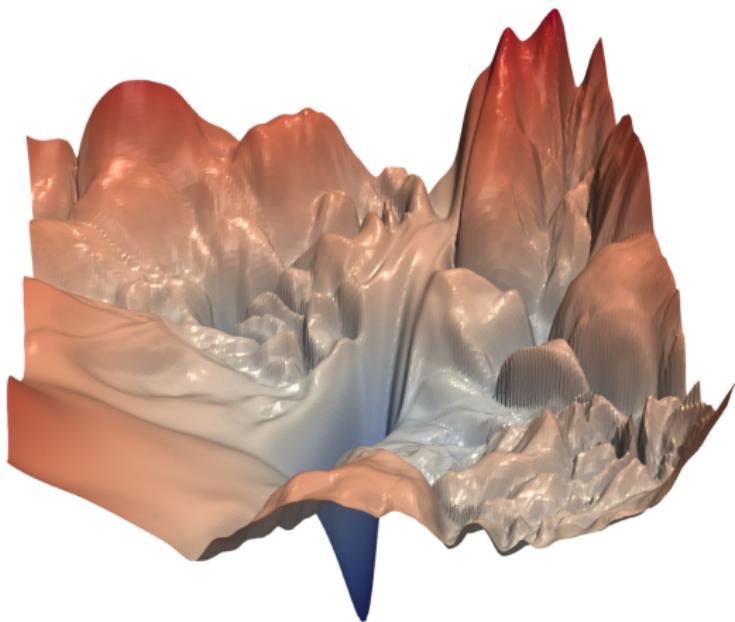
When  $X$ 's dimension reaches the millions, the inverse is essentially intractable.

But we do not need  $\hat{f}_n$ , a close  $\tilde{f}_n$  is good enough for decision making.

Therefore, instead of solving for the best parameters, we just need to approximate it well enough.

# Statistical Learning Theory

## Loss Landscape for Classical Neural Networks



Hao Li, et.al. Visualizing the Loss Landscape of Neural Nets. NeurIPS 2018.

# Statistical Learning Theory

## Gradient Descent

### Idea:

- Given any starting parameters, the gradient indicates the direction of local maximal change.
- If we obtain new parameters by moving old parameter along its gradient, the new ones will give smaller loss (if we are careful).
- We can repeat this procedure until we are happy with the result.

# Statistical Learning Theory

## Directional Gradient Recap

- We say a function  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  is differentiable at  $x \in \mathbb{R}^n$  if

$$\lim_{v \rightarrow 0} \frac{f(x + v) - f(x) - g^T v}{\|v\|_2} = 0,$$

for some gradient vector  $g \in \mathbb{R}^n$  and displacement vector  $v \in \mathbb{R}^n$

- If it exists, this  $g$  is unique and is called the gradient of  $f$  at  $x$  with notation

$$g = \nabla f(x)$$

- It can be shown that

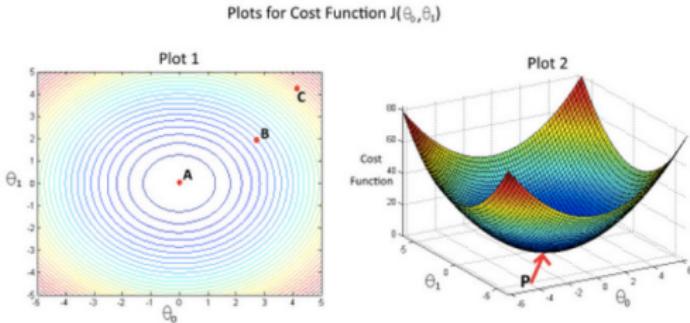
$$\nabla f(x) = \begin{pmatrix} \partial_{x_1} f(x) \\ \vdots \\ \partial_{x_n} f(x) \end{pmatrix}$$

# Statistical Learning Theory

## Contour Graphs

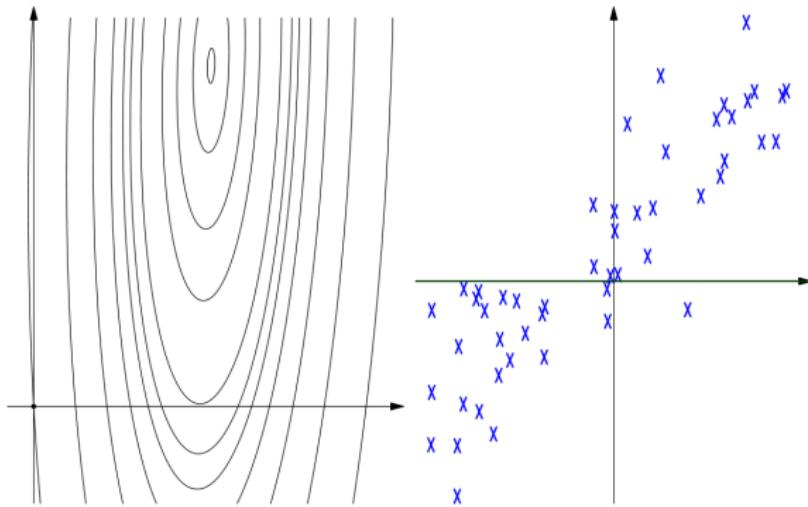
Imagine we are solving a simple linear regression problem:  
 $y = \theta_0 + \theta_1 x$  with loss function:

$$J(\theta_0, \theta_1) = \sum_n (y_i - (\theta_0 + \theta_1 x_i))^2$$




# Statistical Learning Theory

## Negative Gradient Steps



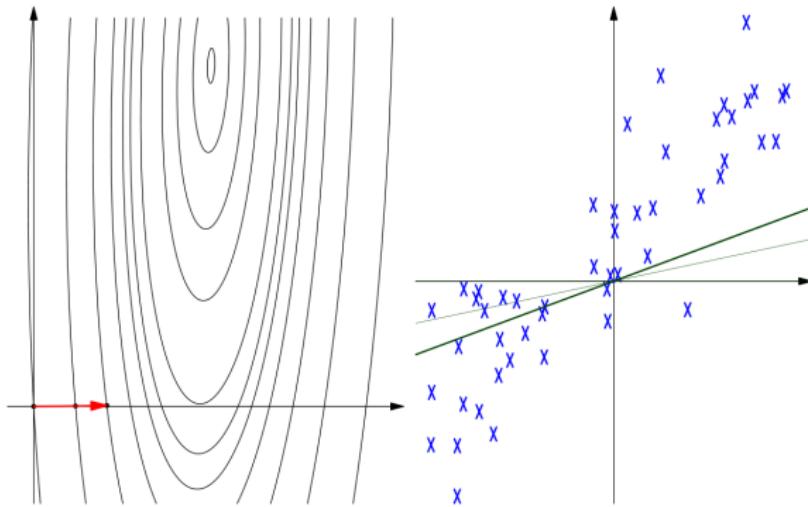
# Statistical Learning Theory

## Negative Gradient Steps




# Statistical Learning Theory

## Negative Gradient Steps



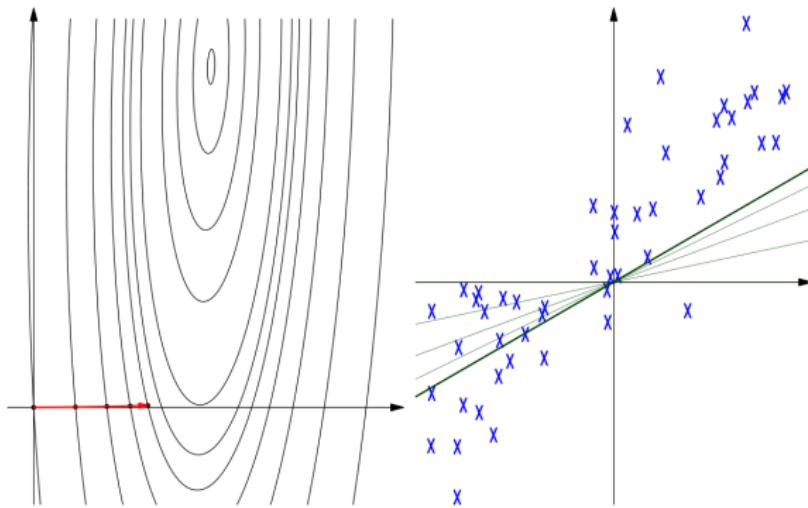
# Statistical Learning Theory

## Negative Gradient Steps




# Statistical Learning Theory

## Negative Gradient Steps



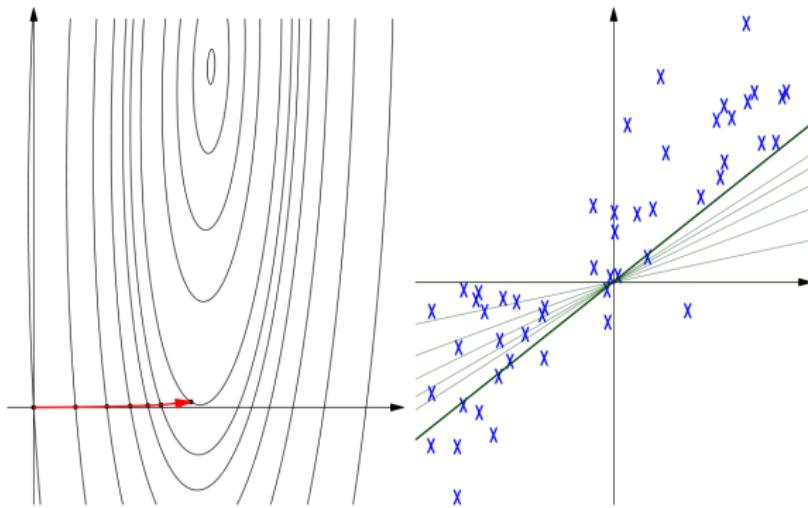
# Statistical Learning Theory

## Negative Gradient Steps



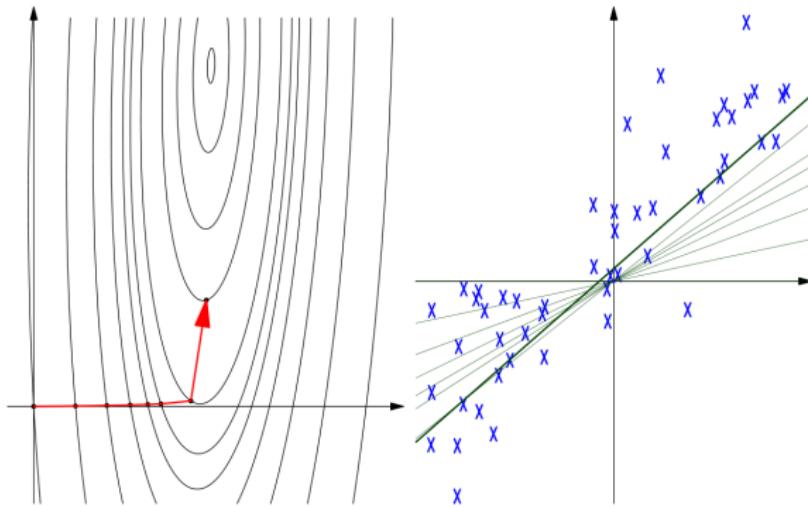
# Statistical Learning Theory

## Negative Gradient Steps



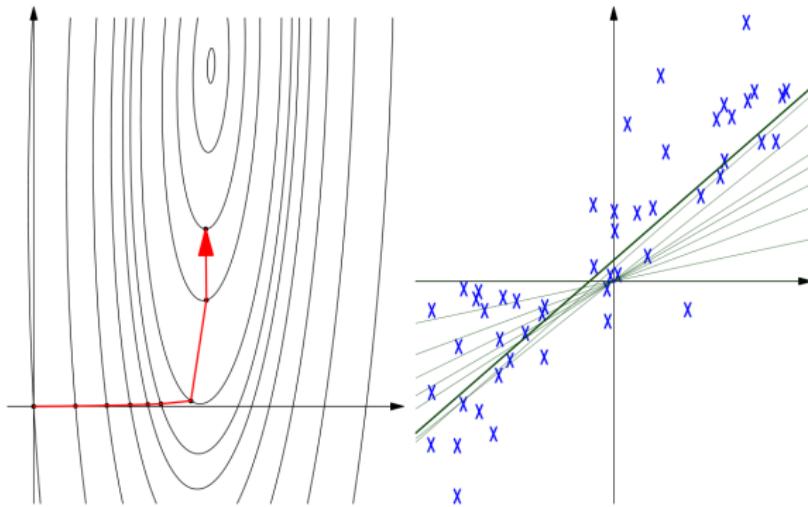
# Statistical Learning Theory

## Negative Gradient Steps



# Statistical Learning Theory

## Negative Gradient Steps



# Gradient Descent

## Gradient Descent

**Goal:** find  $\theta^* = \arg \min_{\theta} J(\theta)$

$\theta_0 :=$  [initial condition] (can be randomly chosen)

$i := 0$

**while** not [termination condition] **do**

    compute  $\nabla J(\theta_i)$

$\alpha :=$  [choose learning rate at iteration  $i$ ]

$\theta_{i+1} := \theta_i - \alpha \nabla J(\theta_i)$

$i := i + 1$

**end while**

**return**  $\theta_i$

# Gradient Descent

## First-order Gradient Descent

**explain how to take gradient/calculus etc. It will be present in next class Objective:** only calculating and using gradient information  $\nabla_{\theta}L(\theta)$  on deciding where to move on loss surface.

**Pros:** Computational Efficiency, Low Memory Footprint, More Scalable

**Cons:** Slow Convergence, Struggle with Curvature, Can stuck at Saddle Points

**Example Algorithms:** SGD, SGD + momentum, Adam, AdamW, RMSProp

# Gradient Descent

## Second-order Gradient Descent

**Objective:** calculating and using both gradient information  $\nabla_{\theta}L(\theta)$  and curvature information  $\nabla_{\theta}^2L(\theta)$  (i.e. Hessian approx) on deciding where to move on loss surface.  
(e.g.  $\theta_{i+1} = \theta_i - \alpha(\nabla^2L(\theta_i))^{-1}\nabla L(\theta_i)$ )

**Pros:** Fast Convergence, Escapes Saddle Points, Curvature awareness

**Cons:** Computational Expense (inverting the Hessian matrix is  $O(n^3)$ ), Memory Explosion (storing full hessian matrix)

**Example Algorithms:** Newton's method, Newton Schulz, Quasi-Newton (L-BFGS), Muon

# Things to review

## Calculus

- Gradients, taking (partial) derivatives

## Linear Algebra

- Matrix computation, matrix derivatives
- Example: compute  $\frac{\partial x^T A x}{\partial x}$ , where  $A$  is a matrix and  $x$  is a vector

Suggested textbooks on Course Website  
[\(https://nyu-dsga-1003.github.io/sp26/\)](https://nyu-dsga-1003.github.io/sp26/)